



HDU-ITMO Joint Institute
杭州电子科技大学 圣光机联合学院

Computer Systems Design

Laboratory work 2

Hardware optimization to PPA constraints

Author:

Alexander Antonov

Associate Professor

Lecturer:

Ivan Lukashov

lukashov@itmo.ru



Contents

1. Objectives	3
2. Overview	3
3. Prerequisites	3
4. Task 3	
5. Guidance	3
6. Variants	13

1. OBJECTIVES

- Students can optimize hardware designs using pipelining
- Students can use Xilinx Vivado Design Suite to gather metrics from implemented design
- Students understand how to achieve PPA trade-off using multi-cycle and pipelined approaches

2. OVERVIEW

Laboratory work 2 covers designing optimized versions of hardware units according to various PPA (performance, power, area) constraints. Students have to re-design previous version of their hardware from Lab 1 according to pipelined approach. Pipelining ubiquitously used in real-world designs – in particular, almost all high-performance designs (CPUs, accelerators, interconnects, memory hierarchy components, etc.) have pipelined organization.

3. PREREQUISITES

Prerequisites are the same as for Lab 1:

1. Xilinx Vivado 2019.1 HLx Edition (free for target board, available at <https://www.xilinx.com/support/download.html>).
2. ActiveCore baseline distribution (available at <https://github.com/AntonovAlexander/activecore>)
3. (for FPGA prototyping) Digilent Nexys 4 DDR FPGA board (<https://store.digilentinc.com/nexys-4-ddr-artix-7-fpga-trainer-board-recommended-for-eee-curriculum/>)
4. (for FPGA prototyping) working Python 3 installation with `pyserial` package

4. TASK

1. Design pipelined implementation in synthesizable SystemVerilog HDL according to your variant
2. Integrate your design with UDM bus master module
3. Write the testbench and simulate to verify correctness of your design
4. Implement your design, gather and analyze metrics of the implementation
5. (if FPGA board available) Perform HW testing on FPGA board
6. (optional) Evaluate top achievable frequency for the designed implementations
7. (optional) Evaluate power consumption of your designs using Vivado power analysis tool
8. Package your solution and submit to the teacher's email

5. GUIDANCE

Detailed guidance will be provided using the example of a module with the same functionality as in Lab 1.

1. Specify the schedule and microarchitectural diagram for your implementation

Pipelined implementation requires the computation to be divided into **stages** that execute in **overlapped** fashion.

Now we develop the **schedule** for our pipelined implementation:

c-step number	operation
0	compare elements in pairs: 0-1; 2-3; 4-5; 6-7; 8-9; 10-11; 12-13; 14-15
1	compare pairing results from stage 0 in pairs: 0-1; 2-3; 4-5; 6-7
2	compare pairing results from stage 1 in pairs: 0-1; 2-3
3	compare pairing results from stage 2

Table 1 Schedule for pipelined implementation

NOTE: Your variants include information about the stages recommended to allocate.

Now we develop microarchitectural diagram defining structure of the module in terms of combinational clouds, registers, and memories:

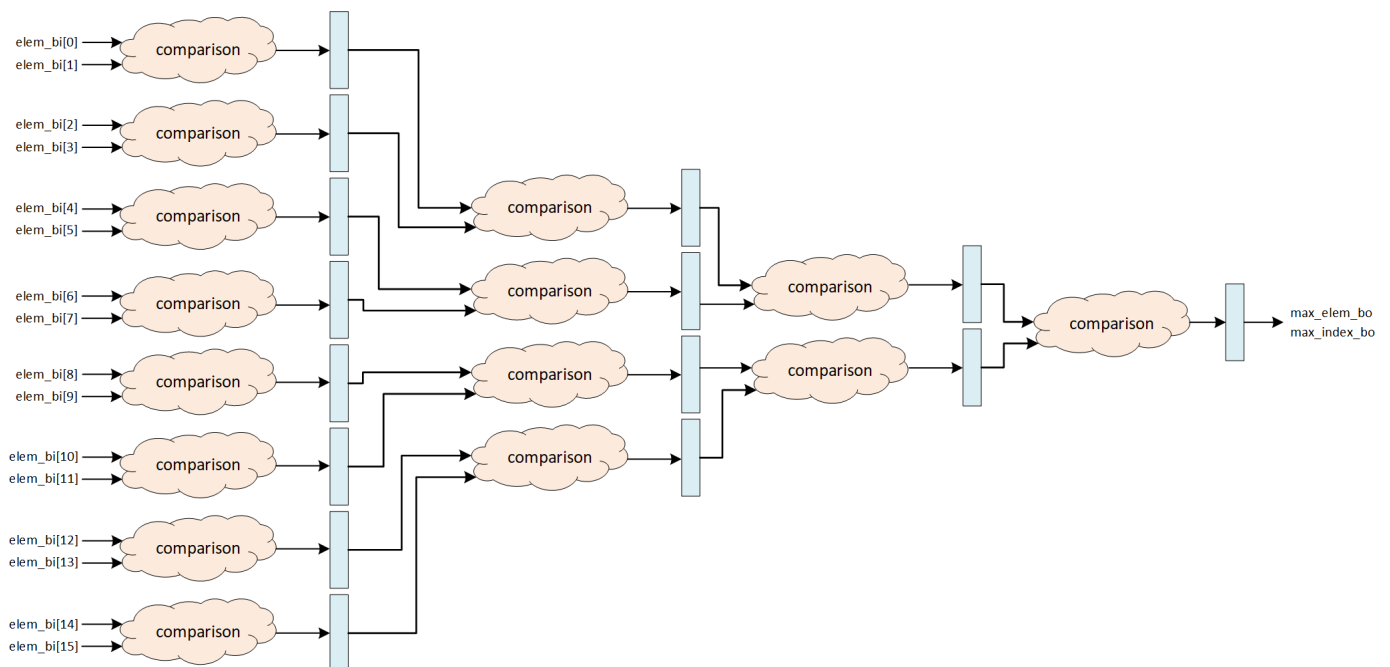


Figure 1 Microarchitectural diagram for fully pipelined implementation

2. Design pipelined implementation in synthesizable SystemVerilog HDL according to the specified schedule and microarchitectural diagram

Source code for the example module is shown in Listing 1.

```

module FindMaxVal_pipelined (
    input clk_i
    , input rst_i

    , input [31:0] elem_bi [15:0]

    , output logic [31:0] max_elem_bo

```

```

    , output logic [3:0] max_index_bo
);

//// stage 0 ////

// intermediate signals declaration
logic [31:0] max_elem_stage0 [7:0];
logic [31:0] max_index_stage0 [7:0];
logic [31:0] max_elem_stage0_next [7:0];
logic [31:0] max_index_stage0_next [7:0];

// combinational logic
always @*
    begin
        for(integer i=0; i<8; i++)
            begin
                max_elem_stage0_next[i] = 0;
                max_index_stage0_next[i] = 0;
                if (elem_bi[(i<<1)] > elem_bi[(i<<1)+1])
                    begin
                        max_elem_stage0_next[i] = elem_bi[(i<<1)];
                        max_index_stage0_next[i] = i<<1;
                    end
                else
                    begin
                        max_elem_stage0_next[i] = elem_bi[(i<<1)+1];
                        max_index_stage0_next[i] = (i<<1)+1;
                    end
            end
        end

// writing to registers
always @(posedge clk_i)
    begin
        if (rst_i)
            begin
                for (integer i=0; i<8; i++) max_elem_stage0[i] <= 0;
                for (integer i=0; i<8; i++) max_index_stage0[i] <= 0;
            end
        else
            begin
                for (integer i=0; i<8; i++) max_elem_stage0[i] <= max_elem_stage0_next[i];
                for (integer i=0; i<8; i++) max_index_stage0[i] <= max_index_stage0_next[i];
            end
        end

//// stage 1 ////

// intermediate signals declaration
logic [31:0] max_elem_stage1 [3:0];
logic [31:0] max_index_stage1 [3:0];
logic [31:0] max_elem_stage1_next [3:0];
logic [31:0] max_index_stage1_next [3:0];

// combinational logic
always @*
    begin
        for(integer i=0; i<4; i++)

```

```

begin
  max_elem_stage1_next[i] = 0;
  max_index_stage1_next[i] = 0;
  if (max_elem_stage0[(i<<1)] > max_elem_stage0[(i<<1)+1])
    begin
      max_elem_stage1_next[i] = max_elem_stage0[(i<<1)];
      max_index_stage1_next[i] = max_index_stage0[(i<<1)];
    end
  else
    begin
      max_elem_stage1_next[i] = max_elem_stage0[(i<<1)+1];
      max_index_stage1_next[i] = max_index_stage0[(i<<1)+1];
    end
  end
end

// writing to registers
always @(posedge clk_i)
begin
  if (rst_i)
    begin
      for (integer i=0; i<4; i++) max_elem_stage1[i] <= 0;
      for (integer i=0; i<4; i++) max_index_stage1[i] <= 0;
    end
  else
    begin
      for (integer i=0; i<4; i++) max_elem_stage1[i] <= max_elem_stage1_next[i];
      for (integer i=0; i<4; i++) max_index_stage1[i] <= max_index_stage1_next[i];
    end
end

///// stage 2 /////

// intermediate signals declaration
logic [31:0] max_elem_stage2 [1:0];
logic [31:0] max_index_stage2 [1:0];
logic [31:0] max_elem_stage2_next [1:0];
logic [31:0] max_index_stage2_next [1:0];

// combinational logic
always @*
begin
  for(integer i=0; i<2; i++)
    begin
      max_elem_stage2_next[i] = 0;
      max_index_stage2_next[i] = 0;
      if (max_elem_stage1[(i<<1)] > max_elem_stage1[(i<<1)+1])
        begin
          max_elem_stage2_next[i] = max_elem_stage1[(i<<1)];
          max_index_stage2_next[i] = max_index_stage1[(i<<1)];
        end
      else
        begin
          max_elem_stage2_next[i] = max_elem_stage1[(i<<1)+1];
          max_index_stage2_next[i] = max_index_stage1[(i<<1)+1];
        end
    end
end
end

```

```

// writing to registers
always @(posedge clk_i)
begin
  if (rst_i)
    begin
      for (integer i=0; i<2; i++) max_elem_stage2[i] <= 0;
      for (integer i=0; i<2; i++) max_index_stage2[i] <= 0;
    end
  else
    begin
      for (integer i=0; i<2; i++) max_elem_stage2[i] <= max_elem_stage2_next[i];
      for (integer i=0; i<2; i++) max_index_stage2[i] <= max_index_stage2_next[i];
    end
end

///// stage 3 /////

// intermediate signals declaration
logic [31:0] max_elem_next;
logic [3:0] max_index_next;

// combinational logic
always @*
begin
  max_elem_next = 0;
  max_index_next = 0;
  if (max_elem_stage2[0] > max_elem_stage2[1])
    begin
      max_elem_next = max_elem_stage2[0];
      max_index_next = max_index_stage2[0];
    end
  else
    begin
      max_elem_next = max_elem_stage2[1];
      max_index_next = max_index_stage2[1];
    end
end

// writing to registers
always @(posedge clk_i)
begin
  if (rst_i)
    begin
      max_elem_bo <= 0;
      max_index_bo <= 0;
    end
  else
    begin
      max_elem_bo <= max_elem_next;
      max_index_bo <= max_index_next;
    end
end

endmodule

```

Listing 1 **Source code of FindMaxVal_pipeline module in SystemVerilog HDL**

3. Integrate the custom design with UDM bus master module

The CSR interface is the same as for combinational implementation.

NOTE: don't forget to connect clock and reset signal to `clk_gen` and `srst` signals respectively (changed lines are highlighted in cyan).

```

module NEXYS4_DDR
#( parameter SIM = "NO" )
(
    input    CLK100MHZ
    , input  CPU_RESETN

    , input  [15:0] SW
    , output logic [15:0] LED

    , input  UART_TXD_IN
    , output UART_RXD_OUT
);

localparam UDM_BUS_TIMEOUT = (SIM == "YES") ? 100 : (1024*1024*100);
localparam UDM_RTX_EXTERNAL_OVERRIDE = (SIM == "YES") ? "YES" : "NO";

logic clk_gen;
logic pll_locked;

sys_clk sys_clk
(
    .clk_in1(CLK100MHZ)
    , .reset(!CPU_RESETN)
    , .clk_out1(clk_gen)
    , .locked(pll_locked)
);

logic arst;
assign arst = !(CPU_RESETN & pll_locked);

logic srst;
reset_cntrl reset_cntrl
(
    .clk_i(clk_gen),
    .arst_i(arst),
    .srst_o(srst)
);

logic udm_reset;

MemSplit32 udm_bus();

udm
#(
    .BUS_TIMEOUT(UDM_BUS_TIMEOUT)
    , .RTX_EXTERNAL_OVERRIDE(UDM_RTX_EXTERNAL_OVERRIDE)
) udm (
    .clk_i(clk_gen)
    , .rst_i(srst)

    , .rx_i(UART_TXD_IN)
    , .tx_o(UART_RXD_OUT)

```



```

    , .rst_o(udm_reset)

    , .bus_req_o(udm_bus.req)
    , .bus_we_o(udm_bus.we)
    , .bus_addr_bo(udm_bus.addr)
    , .bus_be_bo(udm_bus.be)
    , .bus_wdata_bo(udm_bus.wdata)
    , .bus_ack_i(udm_bus.ack)
    , .bus_resp_i(udm_bus.resp)
    , .bus_rdata_bi(udm_bus.rdata)
);

localparam CSR_LED_ADDR          = 32'h00000000;
localparam CSR_SW_ADDR           = 32'h00000004;
localparam TESTMEM_ADDR          = 32'h80000000;

localparam TESTMEM_WSIZE_POW     = 10;
localparam TESTMEM_WSIZE         = 2**TESTMEM_WSIZE_POW;

logic testmem_udm_enb;
assign testmem_udm_enb = (!(udm_addr < TESTMEM_ADDR) && (udm_addr < (TESTMEM_ADDR +
(TESTMEM_WSIZE*4))));

logic testmem_udm_we;
logic [TESTMEM_WSIZE_POW-1:0] testmem_udm_addr;
logic [31:0] testmem_udm_wdata;
logic [31:0] testmem_udm_rdata;

logic testmem_p1_we;
logic [TESTMEM_WSIZE_POW-1:0] testmem_p1_addr;
logic [31:0] testmem_p1_wdata;
logic [31:0] testmem_p1_rdata;

// testmem's port1 is inactive
assign testmem_p1_we = 1'b0;
assign testmem_p1_addr = 0;
assign testmem_p1_wdata = 0;

ram_dual #(
    .init_type("none")
    , .init_data("nodata.hex")
    , .dat_width(32)
    , .adr_width(TESTMEM_WSIZE_POW)
    , .mem_size(TESTMEM_WSIZE)
) testmem (
    .clk(clk_gen)

    , .dat0_i(testmem_udm_wdata)
    , .adr0_i(testmem_udm_addr)
    , .we0_i(testmem_udm_we)
    , .dat0_o(testmem_udm_rdata)

    , .dat1_i(testmem_p1_wdata)
    , .adr1_i(testmem_p1_addr)
    , .we1_i(testmem_p1_we)
    , .dat1_o(testmem_p1_rdata)
);

```

```

assign udm_bus.ack = udm_bus.req;    // bus always ready to accept request
logic csr_resp, testmem_resp, testmem_resp_dly;
logic [31:0] csr_rdata;

// CSR instantiation
logic [31:0] csr_elem_in [15:0];
logic [31:0] csr_max_elem_out;
logic [3:0] csr_max_index_out;

// module instantiation
FindMaxVal_pipelined FindMaxVal_inst (
    .clk_i(clk_gen)
    , .rst_i(srst)

    , .elem_bi(csr_elem_in)

    , .max_elem_bo(csr_max_elem_out)
    , .max_index_bo(csr_max_index_out)
);

// bus request
always @(posedge clk_gen)
begin

    testmem_udm_we <= 1'b0;
    testmem_udm_addr <= 0;
    testmem_udm_wdata <= 0;

    csr_resp <= 1'b0;
    testmem_resp_dly <= 1'b0;
    testmem_resp <= testmem_resp_dly;

    if (srst) LED <= 16'hffff;

    if (srst)          // asserting default values to input CSRs on reset
    begin
        for (int i=0; i<16; i++)
        begin
            csr_elem_in[i] <= 0;
        end
    end

    if (udm_bus.req && udm_bus.ack)
    begin

        if (udm_bus.we)    // writing
        begin
            if (udm_bus.addr == CSR_LED_ADDR) LED <= udm_wdata;
            if (udm_bus.addr[31:28] == 4'h1)    csr_elem_in[udm_bus.addr[5:2]] <=
udm_bus.wdata;
            if (testmem_udm_enb)
            begin
                testmem_udm_we <= 1'b1;
                testmem_udm_addr <= udm_addr[31:2];    // 4-byte aligned access only
                testmem_udm_wdata <= udm_wdata;
            end
        end
    end

```

```

else          // reading
begin
  if (udm_bus.addr == CSR_LED_ADDR)
    begin
      csr_resp <= 1'b1;
      csr_rdata <= LED;
    end
  if (udm_bus.addr == CSR_SW_ADDR)
    begin
      csr_resp <= 1'b1;
      csr_rdata <= SW;
    end
  if (udm_bus.addr == 32'h20000000)
    begin
      csr_resp <= 1'b1;
      csr_rdata <= csr_max_elem_out;
    end
  if (udm_bus.addr == 32'h20000004)
    begin
      csr_resp <= 1'b1;
      csr_rdata <= csr_max_index_out;
    end
  if (testmem_udm_enb)
    begin
      testmem_udm_we <= 1'b0;
      testmem_udm_addr <= udm_addr[31:2];      // 4-byte aligned access only
      testmem_udm_wdata <= udm_wdata;
      testmem_resp_dly <= 1'b1;
    end
end
end
end

// bus response
always @*
begin
  udm_bus.resp = csr_resp | testmem_resp;
  udm_bus.rdata = 0;
  if (csr_resp) udm_bus.rdata = csr_rdata;
  if (testmem_resp) udm_bus.rdata = testmem_udm_rdata;
end


endmodule

```

Listing 2 **Source code of the updated NEXYS4_DDR.sv module**

4. Write the testbench and simulate to verify correctness of the design

The test procedure is the same as for multi-cycle implementation. Waveform for the simulation is shown in Figure 2.



In the waveform, stage-by-stage propagation of data through the pipelined datapath can be observed. The simulation is correct, the module works as intended.

NOTE: In correct linear pipeline, each stage should read data only from the previous stage. Accordingly, a single input data change should cause only a single output data change after several clock cycles (equal to pipeline length). In case a single input data causes the output to change several times, double-check that each stage reads data only from the previous stage. Make copies of pipeline registers for all data needed for computing on corresponding stages.

5. Implement the design, gather and analyze metrics of the implementation

Press “Generate Bitstream” to run implementation and obtain the image for FPGA device.

Metric values are the following:

- Timing:
 - WNS: 4.883 ns (fine)
 - TNS: 0 ns (fine)
- Performance:
 - Clock frequency: 10 ns (100 MHz)
 - Initiation Interval: 1 clock cycle; 10 ns
 - Throughput: 1 op/cycle; 100 Mop/second
 - Latency: 4 clock cycles (equal to schedule length); 40 ns
- HW resources (Implementation → Open Implemented Design → Report Utilization):
 - LUTs: 498
 - FFs (registers): 506

6. (if FPGA board available) Perform HW testing on FPGA board

Python tests and HW testing procedure are the same as in Lab 1.



7. (optional) Evaluate top achievable frequency for the designed implementations

Increase output frequency of `sys_clk` PLL until timing starts failing for each designed implementation. Evaluate top performance of each implementation based on top frequency value when timing closure is achieved.

8. (optional) Evaluate power consumption of your designs using Vivado power analysis tool

Vivado power analysis tool is launched from “Implementation” → “Open implemented design” → “Report power”.

9. Package your solution and submit to the teacher’s email

The package content is equal to Lab 1, but should also include:

- schedule of designed pipeline;
- microarchitectural diagram.

6. VARIANTS

Same as for Lab 1.